

JOÃO MATEUS DE FREITAS VENEROSO

**WEB DATA EXTRACTION IN
SEMI-STRUCTURED DOCUMENTS**

Dissertation proposal presented to the
Graduate Program in Computer Science of
the Federal University of Minas Gerais in
partial fulfillment of the requirements for
the degree of Master in Computer Science.

ADVISOR: BERTHIER RIBEIRO DE ARAÚJO NETO

Belo Horizonte

November 2017

Contents

1	Introduction	1
2	Related Work	3
3	Methodology	6
3.1	Crawler	6
3.2	Data extraction	7
3.2.1	Proposed solution	8
3.2.2	Experiments	12
3.2.3	Further Research	13
4	Schedule	14
	Bibliography	15

Chapter 1

Introduction

Web data extraction is the task of automatically extracting structured information from unstructured or semi-structured web documents. It is a subset of the broader field of Information Extraction, and thus it faces many of the same challenges.

Structured information such as that found in a well organized relational database must conform to an underlying data model, namely an abstract model that formalizes the entities and relationships in a given application domain. Unstructured information however is not organized according to a logical model, therefore useful bits of data won't be arranged cohesively and will occasionally be permeated by chunks of irrelevant information.

Typically, Information Extraction tasks consist of mapping unstructured or poorly structured data to a semantically well defined structure. The input is usually composed of a set of documents that describe a group of entities in a similar manner, while the Information Extraction task deals with identifying those entities and organizing them according to a template.

As an example, consider a collection of novels and the task of identifying the name of the main character in each novel. For this task, the model must first identify proper nouns and then understand the text sufficiently to allow inference on the relative importance of each noun.

To achieve such a goal it is often useful to employ methods developed in the disciplines of Information Retrieval and Natural Language Processing. The former has achieved a great deal of success in the task of classifying documents according to statistical properties and the latter led to huge improvements in modelling human language. Many times, the various methods employed in these disciplines lead to different approaches in the field of Information Extraction.

In the scope of this work, we are interested in the semi-structured data usu-

ally found in HTML web documents. Web documents most often lie in between the structured-unstructured data paradigm, meaning that they take a rather relaxed approach in regard to formal structure. Hierarchy, element disposition, class names, and other features related to document structure and indirectly associated with the data itself are valuable information in the task of identifying entities and determining relationships. So much that many times they are the major source of information for classification purposes, such as when extracting information from standardized tables. However, far from a structured database, web documents usually provide very limited structure to otherwise unstructured data such as that found in free text.

Our current focused research interest regarding web data extraction is collecting computer science researcher information from university websites. We need this information in order to compare the reputations of national and international research groups in the area of computer science with the academic reputation ranking metric proposed by Ribas et al. [2015] based on random walks over reputation graphs.

It showed promising results when ranking publication venues and individual authors in the area of Computer Science by using information publically available in the DBLP repository (<http://dblp.uni-trier.de/>). However the DBLP database has sparse information about author affiliation and multiple records are outdated. To remedy this problem, in the past few years researchers from UFMG have been laboriously collecting this data manually, however this process is tedious and inefficient. Currently, we only have affiliation information for about 1% of the authors, being most of them from USA and Brazil and this is not nearly enough to allow broad international research group comparison.

Up to now, our goal has been to build an automatic information extraction system for collecting author affiliation information from university websites. It has already achieved significant progress which shall be described further in section 3. However, the ideas developed in this concrete case will be further improved in order to construct a more general approach to the broader Information Extraction task.

The research here described hopes to contribute by proposing a novel approach to the main Information Extraction task, making the computer science affiliation database available for further research and testing the quality of academic reputation metrics regarding research groups.

Chapter 2

Related Work

In the last 20 years, the exponential growth of public information in the web has led to the development of a number of different approaches to the problem of web data extraction. Traditionally, the task was solved by designing special purpose programs called wrappers to pinpoint relevant data and store it in some structured format. The multiple tools varied wildly according to their degree of automation.

It was readily perceived that manual wrapper generation was a rather tedious and error prone process, unsuited for large scale operations. Already in 2000, Kushmerick [2000] advocated for wrapper induction, a technique for automatically constructing wrappers. This approach is known in the literature by the acronym WIEN (Wrapper Induction ENvironment).

Web data extraction techniques often require some sort of assistance from human experts to boost accuracy. So the main challenge in the field lies in determining an adequate tradeoff between degree of automation and precision.

In 2002, a survey by Laender et al. [2002a] made a thorough classification of the early approaches with a taxonomy based on their main technology, being them: languages for wrapper development, HTML-aware tools, NLP-based tools, Wrapper Induction Tools, Modeling-based tools and Ontology-based tools. Some noteworthy examples from this era are:

TSIMMIS Hammer et al. [1997] and WebOQL Arocena and Mendelzon [1999], which are special purpose languages for building wrappers.

Road Runner Crescenzi et al. [2001], XWRAP Liu et al. [2000] and W4F Sahuguet and Azavant [1999], which are HTML-aware tools that infer meaningful patterns from the HTML structure.

RAPIER Califf and Mooney [1999], SRV Freitag [1998], WHISK Soderland [1999], which are NLP-based tools.

WIEN ?, Soft Mealy Hsu and Dung [1998] and STALKER ? which are wrapper induction methods.

NoDoSE ? and Debye Laender et al. [2002b], which are semi supervised modeling based tools that require some interaction with the user by means of a graphical user interface.

In 2004, Flesca et al. [2004] developed a taxonomy emphasizing the advantages and drawbacks of web data extraction technologies according to the user viewpoint. In 2006 Chang et al. [2006] complemented the previous surveys with new technologies.

In 2008 Sarawagi [2008] classifies wrappers in the following three types: record-level, page-level and site-level wrappers. Record-level wrappers are only able to process single records, page-level wrappers are capable of extracting data from a single page, and site-level wrappers are able to process the whole webpage structure including subpages and their linking structure.

More recently, surveys by Ferrara et al. [2014] and Schulz et al. [2016] updated the previous surveys and included new approaches.

In 2016 Varlamov and Turdakov [2016], argued that the degree of automation can no longer be the main classification criterion for the data extraction systems because unsupervised methods which were widely considered to be the state of the art when dealing with individual websites performed poorly or were inappropriate on cross site extraction tasks. The authors proposed a classification of methods by the extent of their application. The competing approaches were separated into two groups: methods for individual websites and methods that are applicable to whole application domains.

The first group contains most of the earlier approaches, including the supervised approaches: SRV Freitag [1998], RAPIER Califf and Mooney [1999], WHISK Soderland [1999], WIEN ? SoftMealy Hsu and Dung [1998] and STALKER ?; and the unsupervised approaches: RoadRunner Crescenzi et al. [2001] and EXALG Arasu et al. [2003].

The second group is divided between domain specific methods and domain agnostic methods. Domain specific methods are designed for extracting data about a particular application domain across multiple websites. Our researcher name extractor method that will be further described on section 3 falls in this category. Domain specific methods integrate information about the particular application domain in the course of its development and thus are able to achieve superior performance in comparison to domain agnostic methods.

Domain agnostic methods are the most general extraction methods. They can extract information from any application domain from multiple websites. They pose the hardest challenge because the tool must infer data relevance without any prior training

in that particular application domain. Some examples are: ODE Su and Lochovsky [2009], ObjectRunner Abdessalem et al. [2010], and AMBER Furche et al. [2012]. As previously stated, these approaches tend to yield worse results than domain specific methods, however recent research has been improving domain agnostic methods by large steps.

Chapter 3

Methodology

This section describes briefly a statistical NLP approach to solve the researcher affiliation extraction problem presented in section 1. This problem is actually composed of two parts. The first one involves crawling university websites and finding faculty repositories and the second one is extracting the actual researcher affiliation data from these repositories. We will present here the model developed up to now, and then we will present the further research goals that are still to be reached.

3.1 Crawler

The crawling problem has been partially solved by online aggregators that compiled university internet domains for most of the academic world. One of these aggregators is <https://univ.cc/>, which contains links for the main websites of 9553 universities from 207 countries. These root domains were fed as seeds to our crawler, that downloaded the entire websites and its subdomains, amounting around 2 million documents. Once all webpages were downloaded the task at hand becomes identifying faculty repositories among the downloaded pages.

This task can be accomplished quite efficiently by means of supervised machine learning models once we have labeled a subset of the university webpages. The labeling was manually done for 2440 pages, of which 373 were faculty repositories and 2067 were other types of pages. We took care to select a proportional number of pages from each country.

Since the classification task is not our main research focus, the results will only be presented here briefly. Our classifier employed only 112 features:

- 60 for the number of occurrences of a set of relevant keywords in the URL, the

page title, h1, h2 and h3 tags. Each element having its own set of features.

- 49 for the number of occurrences of a different set of keywords in the rest of the text of the page.
- 1 feature measuring the document length.
- 1 feature for the number of names found by a simplified version of the name extractor.

Model	Accuracy	Standard Deviation
Random Forest	0.93	0.02
Gaussian Naive Bayes	0.88	0.03
Linear SVM	0.85	0.05
Logistic Regression	0.90	0.02

Table 3.1. Classifier



Four models were tested by means of 5 fold cross validation, the results are presented in table 3.1. All models except Support Vector Machines achieved a reasonable performance, however, since the Random Forest Classifier presented a more stable predictive capability, it was chosen as the classifier method for labeling the remaining webpages.

3.2 Data extraction

The data extraction stage consists of extracting researcher names from faculty directories. The main challenge here is constructing a general enough approach that can be used effectively across different domains without extracting too much garbage together with the valuable information. This task faces a typical precision vs. recall tradeoff.


In order to understand the complexity of this extraction task take for example the staff page for the intelligent robotics laboratory of Osaka University shown in figure 1. Say we want to extract the name, position, email and picture of all members. It is easy to see there is some sort of structure to the information we want to extract from this particular website, however, parts of the information are missing, repeated or disposed differently for each particular member. If we want to go further it may be necessary to only extract information from full members, a task that may pose a harder challenge. We may use grouping similarity combined with textual information in order to properly identify the desired entities. In many cases the HTML element's

Graduate School of Engineering Science
Frontier Intelligent System Research Laboratory (CyberAgent, Inc.)

 <p>Specially Appointed Associate Professor Itaru Kuramoto ✉ kuramoto</p>	 <p>Visiting Fellow Jun Baba ✉ Baba</p>
---	---

domain: @irl.sys.es.osaka-u.ac.jp

**Dept. of Adaptive Machine Systems,
Graduate School of Engineering
(Collaborating Division)**

 <p>Professor (Distinguished Professor) Hiroshi Ishiguro ✉ ishiguro</p>

Invited Professor



<p>National Institute of Advanced Industrial Science and Technology domain : @aist.go.jp</p>  <p>Yoshio Matsumoto ✉ yoshio.matsumoto</p>	<p>ATR Intelligent Robotics and Communication laboratories, Human-Robot Interaction Department domain : @atr.jp</p>  <p>Takayuki Kanda ✉ kanda</p>
---	---

Figure 3.1. Example of semi structured information

relative position or CSS class name is sufficient to identify particular occurrences of a same entity. If we do a good enough job at this first task we may extrapolate our model to extract information from similar web pages.

However at this moment we are only interested in extracting researcher names and it is not necessary to identify their status regarding full or partial membership.

3.2.1 Proposed solution

The name extraction problem is no different than a Named Entity Recognition problem, however approaches that typically achieve a high accuracy on free text like Conditional Random Fields, Hidden Markov Models or Maximum Entropy Models perform very poorly in this particular case, because the data we are trying to extract is disposed in a tabular form and the words do not hold dependencies toward each other the same

way they do in long pieces of text. In the name extraction problem, structural HTML features and proper noun conditional probabilities play a key role.

Let $t = (t_1, t_2, \dots, t_n)$ be a sequence of tokens, and $y = (y_1, y_2, \dots, y_n)$ be a sequence of labels where $y_i \in \{N, W\} \forall i \in \mathbb{N}$, such that $y_i = N$ means token i is a name and $y_i = W$ means token i is a word. The problem of extracting names from a sequence of tokens is equivalent to the problem of finding an optimal sequence of labels y^* for a sequence of tokens t :

$$y^* = \underset{y}{\operatorname{argmax}} P(t_i = y_i, t_{i+1} = y_{i+1}, \dots, t_n = y_n) \quad (3.1)$$

where $P(t_i = y_i)$ is the probability that token t_i has label y_i . We may employ the chain rule to explore the relationship between the joint and conditional probabilities. Consider that $P(Y_i) = P(t_i = y_i)$

$$P(Y_1, Y_2, \dots, Y_n) = P(Y_n)P(Y_2|Y_1) \dots P(Y_n|Y_1, Y_2, \dots, Y_{n-1}) \quad (3.2)$$

we could approximate these conditional probabilities the same way we would do for a n-gram language model. However, the conditional probabilities $P(t_i = y_i | t_{i-1}y_{i-1}, \dots)$ are hard to estimate because the joint distribution $P(t_i, y_i)$ depends both on the previous label and the previous token. If we express them in terms of joint probabilities the problem becomes clearer:

$$P(t_i, y_i | t_{i-1}, y_{i-1}, t_{i-2}, y_{i-2}, \dots) \quad (3.3)$$

so we make the assumption that the probability that token t_i has label y_i depends on the values of previous labels but is independent of the previous tokens. For example, given a sequence of tokens $\{John, Hall\}$, the conditional probability $P(Hall | John = name)$ is equivalent to $P(Hall | any\ name)$. In other words, this means that the probability of Hall being a last name is the same regardless of a person's first name, as long as we can make sure that the previous token is a name. The conditional probabilities then become:

$$P(t_i, y_i | t_{i-1}, y_{i-1}, t_{i-2}, y_{i-2}, \dots) = P(t_i, y_i | y_{i-1}, y_{i-2}, \dots) \quad (3.4)$$

Replacing equation 3.5 in equation 3.2, yields:

$$P(Y_1, Y_2, \dots, Y_n) = P(Y_1)P(Y_2|y_1) \dots P(Y_n|y_1, y_2, \dots, y_{n-1}) \quad (3.5)$$

We can once again employ the chain rule of probability to write:

$$P(x_i, y_i | y_1, y_2, \dots, y_{i-1}) = P(y_i | y_1, y_2, \dots, y_{i-1}) P(x_i | y_1, y_2, \dots, y_i) \quad (3.6)$$

Approximating $P(x_i | y_1, y_2, \dots, y_i)$ by $P(x_i | y_i)$ we obtain:

$$P(x_i, y_i | y_1, y_2, \dots, y_{i-1}) = P(y_i | y_1, y_2, \dots, y_{i-1}) P(x_i | y_i) \quad (3.7)$$

Take notice that $P(Y_i)$ is a different notation for $P(x_i = y_i)$ or $P(x_i, y_i)$. So, by replacing equation 3.7 in equation 3.5 we finally obtain:

$$P(Y_1, Y_2, \dots, Y_n) = P(y_1, y_2, \dots, y_n) P(x_1 | y_1) P(x_2 | y_2) \dots P(x_n | y_n) \quad (3.8)$$

Equation 3.8 can be split into two parts: the prior, given by the first part of the equation on the right side and the conditional token probabilities, given by the rest of the equation on the right side.

3.2.1.1 Prior probabilities

In order to obtain the prior probability, we need to acquire estimates for all possible sequences of labels. Considering that label y_i must be either a name or a word, then there are 2^k different combinations for a window of size k .

Let N be a name label and W be a word label, then for a window of size k we would need to estimate prior probabilities for all 2^k possible sequences of labels. In practice, a window of size 4 seems to be accurate enough. In this case, the sequences would be: $\{W, W, W, W\}, \{W, W, W, N\}, \{W, W, N, W\}, \dots$

When names occur next to each other we have no way to tell where the first name ends and the second one starts. In order to delimit name boundaries we need to estimate priors for different sequences of name labels in addition to our previous priors. Let the first and second name labels be N_1 and N_2 , respectively. Then, we need to estimate priors for the sequences: $\{W, N_1, N_1, N_2\}, \{W, N_1, N_1, N_2\}, \{W, N_1, N_2, N_2\}, \dots$. In practice, we are never interested in isolated occurrences of name labels so we can exclude combinations such as $\{W_1, N_1, N_2, N_2\}$.

Most of the times when names happen inside a list they tend to be contained inside a single HTML element. Eventhough this is not always the case, this knowledge can be incorporated as an additional piece of evidence in our model. This evidence becomes specially useful when we are trying to delimit name boundaries. Let $*$ indicate a breaking point in a sequence of labels, so $W, W*, N, N$ means that the tokens taking labels WW are contained inside a single HTML element, while the remaining tokens are inside different HTML elements. We could estimate sequences with multiple breaking

points, however a single breaking point has shown good results. For our window of size 4, we need to estimate all prior probabilities with the 4 possible breaking points: $\{y_1, y_2, y_3, y_4\}, \{y_1^*, y_2, y_3, y_4\}, \{y_1, y_2^*, y_3, y_4\}, \{y_1, y_2, y_3^*, y_4\}$.

3.2.1.2 Conditional token probabilities

We need to estimate conditional token probabilities for both labels: names and words. So we need to know $P(t_i|N)$, the probability that a name is t_i and $P(t_i|W)$, the probability that a word is t_i .

For our experiments, the conditional token probabilities were obtained by maximum likelihood estimation with Laplace smoothing to account for tokens that didn't occur in the corpus. The $P(t_i|N)$ probabilities were estimated over a collection of approximately 1.5 million names from the DBLP database. The $P(t_i|W)$ probabilities were estimated over a corpus of 100 thousand documents obtained during the crawling stage. In the latter case all capitalized words were ignored when estimating probabilities in order to remove most names from the corpus.

Token probabilities can be made more precise by incorporating features in equation 3.8. We do that by changing the token conditional probabilities to:

$$P(t_i, f_1, f_2, \dots, f_n|y_i) = P(t_i|y_i)P(f_1|y_i) \dots P(f_n|y_i) \quad (3.9)$$

where f_i are features, which are assumed to be independent between themselves. The features can be textual or structural. Textual features take textual clues like previous words and token length to predict if a given token is a name. Structural features infer token probabilities based on the HTML structure like tag names and nesting depth. In practice previous and next words weren't particularly effective in empirical tests, possibly due to the fact that most names occurred in lists in the test collection. So they were removed from the features list.

Structural features estimated over the entire corpus end up being too general so they help little in increasing token probability estimates. HTML structure varies radically between different documents such that the only stable characteristic is that names tend to have similar structural contexts in the same faculty directory. For example, if all names appear inside a `<tr>` tag in a given document it does not mean that names tend to appear inside `<tr>` rather than any other tag in other documents. However for that particular document we may be able to identify other names and exclude words by knowing that tokens occurring inside `<tr>` tags have a higher probability of being names.

If our basic algorithm (without structural features) was able to extract a good number of names from a page on the first passing, we may use their structural context to estimate probabilities for our structural features. Of course those tokens that were tagged as words can be used to estimate the word conditional probabilities. On a second passing we can incorporate these improved estimates to boost the model's precision considerably. This process can be repeated multiple times to further increase performance.

Id	Feature	Description
1	Token incidence	How often a token happens in a document's text
2	Token length	The token's character length
3	First+Second parents	First and second HTML parents combined
4	Third parent	Third HTML parent
5	CSS class name	Innermost CSS class valid for the token
6	Child number	The child number in relation to the HTML parent
7	Nesting depth	The number of HTML parents up to root

Table 3.2. Features

Table 3.2 describes the features used in our experiments. A couple of other features were tested, but they weren't included in our analysis because they didn't have a significant impact on the experimental results. Some of them are: next word indicates some location (street, avenue, etc.), previous word is an honorific, token is capitalized, token is the first element in an HTML child, etc.

3.2.2 Experiments

Feature	Precision	Recall	F-measure
None	0.1	0.3	0.3
1	0.1	0.3	0.3
2	0.1	0.3	0.3
1+2	0.1	0.3	0.3

Table 3.3. Textual features experiment

Experiments were made to determine the best model and the best set of features. The test collection was a set of 310 manually labeled faculty directory pages. For each model and each group of features the precision, recall and f-measure were calculated. Extracted names were only considered to be correct when they resulted on an exact match with the test data. All measures were obtained by the averaged results of a 5

fold cross validation run. We first compared the textual features on our base model (NE), which passes only one time over the token list. The results are presented on table 3.3.

Feature	Precision	Recall	F-measure
None	0.1	0.3	0.3
3	0.1	0.3	0.3
4	0.1	0.3	0.3
5	0.1	0.3	0.3
6	0.1	0.3	0.3
4	0.1	0.3	0.3
All	0.1	0.3	0.3

Table 3.4. Structural features experiment

Next, we tested the structural features in a model that passes two times over the token list (NE2) estimating the structural feature probabilities on the second passing. We used the best set of textual features found in the previous experiment. The results are presented on table 3.4.

Number of passings	Precision	Recall	F-measure
1	0.1	0.3	0.3
2	0.1	0.3	0.3
3	0.1	0.3	0.3
4	0.1	0.3	0.3
5	0.1	0.3	0.3

Table 3.5. Multiple passings experiment

Finally we compared models with 1, 2, 3, 4 and 5 passings over the token list using the best set of features from the previous experiments. The results are presented on table 3.5.

3.2.3 Further Research

Our model is finely tuned for the particular problem of researcher name extraction, but we believe this result can be generalized to solve other data extraction problems. Furthermore, the researcher affiliation data collected by the name extractor still needs to be used for evaluating the reputation metric proposed by Ribas et al. [2015] in the research group classification task.

Chapter 4

Schedule

The dissertation project will be divided into 6 tasks that shall be accomplished along 2018. The tasks are:

1. **Literature review:** review the literature about the topic of research
2. **Data analysis:** collect and structure the relevant data.
3. **Baseline:** recreate some key approaches from the literature.
4. **Model implementation:** implement the general approach of the model described in section 3.
5. **Experiments:** compare the implemented model with the baseline.
6. **Dissertation:** write the dissertation and reviewing the text.

Task	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1	x	x	x									
2	x	x	x									
3				x	x	x						
4				x	x	x						
5							x	x	x			
6							x	x	x	x	x	x

Table 4.1. Schedule 2018

Bibliography

- Abdessalem, T., Cautis, B., and Derouiche, N. (2010). ObjectRunner: lightweight, targeted extraction and querying of structured web data. *Proceedings of the VLDB ...*, 3(2):1585--1588. ISSN 21508097.
- Arasu, A., Garcia-Molina, H., Arasu, A., and Garcia-Molina, H. (2003). Extracting structured data from Web pages. *2003 ACM SIGMOD International Conference on Management of Data*, pages 337 -- 348.
- Arocena, G. O. and Mendelzon, A. O. (1999). WebOQL: Restructuring documents, databases, and webs. *Theory and Practice of Object Systems*, 5(3):127--141. ISSN 10743227.
- Califf, M. E. and Mooney, R. J. (1999). Relational learning of pattern-match rules for information extraction. *Computational Linguistics*, 4:9--15. ISSN 15324435.
- Chang, C.-H., Kayed, M., Girgis, M. R., and Shaalan, K. F. (2006). A Survey of Web Information Extraction Systems. *IEEE Transactions on Knowledge and Data Engineering*, 18(10):1411--1428. ISSN 1041-4347.
- Crescenzi, V., Mecca, G., and Merialdo, P. (2001). Roadrunner: Towards automatic data extraction from large web sites. *Proceedings of the 27th International Conference on Very Large Data Bases*, pages 109--118. ISSN 10477349.
- Doan, A., Naughton, J. F., Ramakrishnan, R., Baid, A., Chai, X., Chen, F., Chen, T., Chu, E., Derosé, P., Gao, B., Gokhale, C., Huang, J., Shen, W., and Vuong, B.-q. (2008). Information Extraction Challenges in Managing Unstructured Data. *ACM SIGMOD Record*, 37(4):14--20. ISSN 0163-5808.
- Ferrara, E., De Meo, P., Fiumara, G., and Baumgartner, R. (2014). Web data extraction, applications and techniques: A survey. *Knowledge-Based Systems*, 70:301--323. ISSN 09507051.

- Figueiredo, L. N. L., de Assis, G. T., and Ferreira, A. A. (2017). DERIN: A data extraction method based on rendering information and n-gram. *Information Processing & Management*, 53(5):1120--1138. ISSN 03064573.
- Fiumara, G. (2007). Automated information extraction from Web sources: A survey. *CEUR Workshop Proceedings*, 312:1--9. ISSN 16130073.
- Flesca, S., Manco, G., Masciari, E., Rende, E., and Tagarelli, A. (2004). Web wrapper induction: a brief survey. *AI Communications*, 17(2):57--61. ISSN 0921-7126.
- Freitag, D. (1998). Information Extraction from HTML: Application of a General Machine Learning Approach. *Proceedings of the Fifteenth National/Tenth Conference on Artificial Intelligence/Innovative Applications of Artificial Intelligence*, pages 517--523.
- Furche, T., Gottlob, G., Grasso, G., Orsi, G., Schallhart, C., and Wang, C. (2012). AMBER: Automatic Supervision for Multi-Attribute Extraction. *arXiv preprint*, 1210(5984):1--22.
- Garfield, E. (1955). Citation Indexes for Science. *Science*, 122:108--111.
- Hammer, J., Mchugh, J., and Garcia-molina, H. (1997). Semistructured Data : The TSIMMIS Experience. *Proceedings of the First East-European Symposium on Advances in Databases and Information Systems*, pages 1--8.
- Hsu, C. N. and Dung, M. T. (1998). Generating finite-state transducers for semi-structured data extraction from the Web. *Information Systems*, 23(8):521--538. ISSN 03064379.
- Kaiser, K. and Miksch, S. (2005). Information Extraction. *Technology*, (May):32.
- Kao, H.-A. and Chen, H.-H. (2010). Comment Extraction from Blog Posts and Its Applications to Opinion Mining. *Proceedings of the 7th International Conference on Language Resources and Evaluation*, pages 1113--1120.
- Khalil, S. and Fakir, M. (2017). RCrawler: An R package for parallel web crawling and scraping. *SoftwareX*, 6:98--106. ISSN 23527110.
- Kushmerick, N. (2000). Wrapper induction: efficiency and expressiveness. *Artificial Intelligence*, 118(1-2):15--68. ISSN 00043702.
- Kushmerick, N. and Kushmerick, N. (2003). Finite-state approaches to Web information extraction. *Lecture Notes in Computer Science*, 2700:77--91. ISSN 03029743.

- Kushmerick, N., Weld, D. S., and Doorenbos, R. (1997). Wrapper induction for information extraction. *Intl Joint Conference on Artificial Intelligence IJCAI*, pages 729--737.
- Laender, A., Ribeiro-Neto, B. A., and S.Teixeria, J. (2002a). A brief survey of web data extraction tools. *ACM SIGMOD Record 31(2)*, pages 84--93.
- Laender, A. H. F., Ribeiro-Neto, B., and da Silva, A. S. (2002b). DEByE - Date extraction by example. *Data Knowl. Eng.*, 40(2):121--154. ISSN 0169-023X.
- Liu, L., Pu, C., and Han, W. (2000). XWRAP: an XML-enabled wrapper construction system for Web information sources. *Proceedings of 16th International Conference on Data Engineering*, pages 611--621. ISSN 1063-6382.
- Ribas, S., Ribeiro-Neto, B., Santos, R., Souza e Silva, E., Ueda, A., and Ziviani, N. (2015). *Random Walks on the Reputation Graph*. ACM Press, New York, New York, USA. ISBN 9781450338332.
- Sahuguet, A. and Azavant, F. (1999). Building light-weight wrappers for legacy Web data-sources using W4F. *Proceedings of the 25th VLDB Conference*, 99:738--741.
- Sarawagi, S. (2008). Information extraction. *Foundations and Trends in Databases*, 1(3):261--377. ISSN 1931-7883.
- Schulz, A., Lässig, J., and Gaedke, M. (2016). Practical web data extraction: Are we there yet? — A short survey. *IEEE/WIC/ACM International Conference on Web Intelligence (WI), 2016*, pages 562----567.
- Soderland, S. (1999). Learning Information Extraction Rules for Semi-Structured and Free Text. *Machine Learning*, 34(1):233--272. ISSN 0885-6125.
- Su, W. and Lochvsky, F. H. (2009). ODE : Ontology-assisted Data Extraction. 1(212).
- Varlamov, M. I. and Turdakov, D. Y. (2016). A survey of methods for the extraction of information from Web resources. *Programming and Computer Software*, 42(5):279--291. ISSN 0361-7688.
- Zhai, Y. and Liu, B. (2005). Web data extraction based on partial tree alignment. *Proceedings of the 14th international conference on World Wide Web - WWW '05*, page 76. ISSN 10414347.