Researcher name extraction from faculty directories

JOÃO MATEUS DE FREITAS VENEROSO

Public databases such as DBLP and Google Scholar contain valuable information about the academic environment that have been incredibly useful helping answering numerous research questions. However, in these databases, only a fraction of the records contain researcher affiliation information, and even in the cases where the information is present it is frequently outdated. The problem can be better or worse depending on the field of study. We propose a method to extract researcher names from faculty directories automatically in order to keep up to date information about researcher affiliation and be general enough to solve any name extraction task. Various web data extraction methods have already been proposed in the literature and they typically lack either in generality, because wrappers trained on a set of examples do not perform well when handling different websites, or precision, because domain agnostic methods perform poorly compared to domain specific methods. Our statistical NLP approach solves the name extraction task with a framework that incorporates both textual and structural features to yield an excellent tradeoff between generality and precision. We conducted experiments over a collection of 310 faculty web pages from multiple universities across the world and obtained 94.37% precision, 97.61% recall and 0.9596 F-measure.

Additional Key Words and Phrases: Information extraction, statistical classifier, natural language processing

Fig. 1. Example of faculty directory
.

1 INTRODUCTION

Web data extraction is the task of automatically extracting structured information from unstructured or semi-structured web documents. Tipically, Information Extraction tasks consist of mapping unstructured or poorly structured data to a semantically well defined structure. The input is most commonly composed of a set of documents that describe a group of entities in a similar manner, while the Information Extraction task deals with identifying those entities and organizing them according to a template.

HTML documents most often lie in between the structured / unstructured data paradigm, which means that authors take a rather relaxed approach in regard to formal structure. Hierarchy, element disposition, class names, and other features related to the document structure and indirectly associated with the data itself are valuable information in the task of identifying entities and determining relationships, yet we cannot expect these features to be completely constrained by any underlying pattern. Like natural language, organization patterns tend to follow some guidelines but are in no way subject to strict rules.

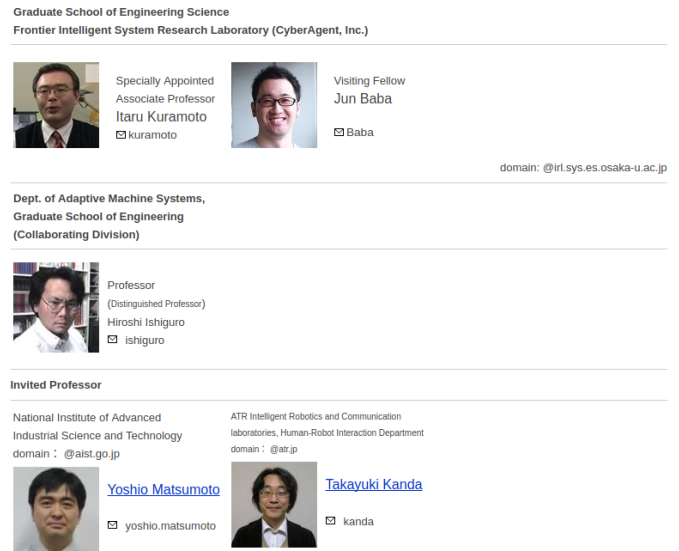We are interested in the task of name extraction, particularly extracting researcher names from university websites to complement data from public databases such as the DBLP repository (http://dblp.uni-trier.de/) in order to allow broad international research group comparison and enquiring about publication patterns. The DBLP database has sparse information about author affiliation and only contains information about computer science researchers. With our probabilistic method we hope to build a general approach to solve this problem without having to rely on metadata from PDF papers and handle the name extraction problem in general.

In order to acknowledge the complexity of this extraction task take for example a snippet of the staff page for the intelligent robotics laboratory from Osaka University shown in figure 1. There is some structure to the way member profiles are arranged, but the organization is rather flexible even considering this single website. Other websites can show very different patterns, ranging from tables and lists to free form.

Researcher names can happen inside plain text, similar to typical named entity recognition scenarios. Names may be part of larger sentences such as in "Michael Johnson Chair" and "John Doe Avenue" yielding false positives. Names can be composed of common words (e.g. Summer Hall) yielding false negatives. In short, there is no one rule that fits all cases.

State-of-the-art named entity recognition approaches such as Conditional Random Fields do not perform so well in information extraction cases such as the one presented in figure 1, because the text is insufficient to provide any contextual information about the semantic category of a word.

Since we lack a database with all possible name combinations, we propose a hollistic statistical method that accounts for discrepancies in data organization by assigning probabilities to sequences of tokens without relying on per-website training. Our method of label assignement resembles a Naive Bayesian classifier without the assumption of token independence. We also rely as little as possible on contextual information to avoid the mistakes of other classifiers. The base model already achieves recall and precision rates above 90%, but the performance can be increased even further by incorporating HTML structural features to estimate better probabilities, specially to avoid filtering out false negatives.

2 RELATED WORK

In the last 20 years, the astonishing growth of public information in the web has led to the development of a number of different approaches to the problem of web data extraction. Traditionally, the task was solved by designing special purpose programs called wrappers to recognize relevant data and store it in a structured format. These early tools varied wildly relative to their degree of automation.

It was readily perceived that manual wrapper generation was a rather tedious and error prone process, unsuited for large scale operations. Wrappers tend to break frequently because they rely heavily on web page features that can change often. So, in the late nineties, several authors advocated for wrapper induction, a technique that consists of automatically constructing wrappers from a small set of examples by identifying delimiters or context tokens that single out the desired attributes. Some remarkable wrapper induction methods are WIEN [2], Soft Mealy [3] and STALKER [4].

Despite being better than constructing wrappers manually, wrapper induction methods still suffered from a lack of expressive power and flexibility. These methods had trouble handling records with missing attributes or unusual structures because patterns could only be identified if they happened at least once in the examples.

Other approaches such as NoDoSE ([5]) and Debye ([6]) brought greater flexibility to wrapper induction methods by requiring a greater level of human interaction through graphical user interfaces. Web data extraction techniques often require some sort of assistance from human experts to boost accuracy. One of the main challenges in the field lies in determining an adequate tradeoff between the degree of automation and the precision and recall of the data extraction tool.

In order to automate the task of web data extraction completely some approaches, such as Road Runner [7], removed entirely the need for data examples. Road Runner parses documents belonging to a same class (e.g. books on Amazon) and generated wrappers based on their similarities and differences, yielding comparable results to those obtained by wrapper induction methods. However like previous approaches, it was unsuited for cross site extraction tasks because the learned rules weren't general enough.

NLP based approaches aimed at extracting more general rules that could possibly be employed over multiple websites. RAPIER [8] is a method of rule extraction that uses information such as part-of-speech tags and semantic classes from a lexicon to derive patterns from a set of training examples. This approach is more

flexible than the wrapper induction methods, however it achieves much lower rates of recall and precision.

In 2002, a survey by Laender et al. [9] made a thorough classification of the early approaches with a taxonomy based on their main technology, being them: languages for wrapper development, HTML-aware tools, NLP-based tools, Wrapper Induction Tools, Modeling-based tools and Ontology-based tools. Some noteworthy examples from this era are:

- TSIMMIS [10] and WebOQL [11], which are special purpose languages for building wrappers.
- Road Runner [7], XWRAP [12] and W4F [13], which are HTML-aware tools that infer meaningful patterns from the HTML structure.
- RAPIER [8], SRV [14], WHISK [15], which are NLP-based tools.
- WIEN [2], Soft Mealy [3] and STALKER [4] which are wrapper induction methods.
- NoDoSE [5] and Debye [6], which are semi supervised modeling based tools that require some interaction with the user by means of a graphical user interface.

In 2006, Chang et. al. [16] complemented the previous surveys with semisupervised technologies such as Thresher [17], IEPAD [18] and OLERA [19]. They differed from supervised and unsupervised methods because they either needed only a rough description of data from users for extraction rule generation or some level of post processing that needed user attention. The survey also mentioned newer unsupervised methods such as DeLa [20], Exalg [21] and Depta [22].

Most of the early information extraction systems were rule-based with either manual rule description or automatic rule learning from examples, thus they suffered from a lack of flexibility when dealing with noisy and unstructured data. Huge progress in the field of statistical learning led to the development of statistical models that tried to solve this problem.

In 2008, Sarawagi [23] produced a survey that classified wrappers in rule-based methods, statistical methods and hybrid models, bringing together the fields of named entity recognition, relationship extraction and information extraction. The rule based methods encompass most of the previous models. The statistical methods convert the extraction task into a token labeling task, identifying the target entities through the assignment of labels. Any classifiers such as a Support Vector Machines, Logistic Classifiers or Neural Networks could be employed to perform this task. However Hidden Markov Models, Maximum Entropy Taggers and Conditional Random Fields tend to perform better at most extraction tasks because of the way they model token and label dependencies. Hybrid models incorporate both rule-based and statistical methods.

Statistical models have proven to be reliable tools for performing numerous NLP tasks. However, information from web documents relevant to data extraction tasks is usually arranged in a tabular form rather than in a plain text format. Therefore, typical state of the art classifiers can yield poor results if they rely exclusively on textual information. In the web information extraction task, the document structure must be incorporated as a feature in an effective classifier.

More recently, surveys by [24] and [25] updated the previous surveys with some interesting innovations. Some examples are: the Visual Box Model [26], a data extraction system that produces a visualization of the web page to exploit visual cues to identify data presented in a tabular form; automatic wrapper adaptation [27], a tecnique that tries to reduce the cost of wrapper maintenance by measuring the similarity of HTML trees and adapting wrappers to the new page structure; AutoRM [28], a method to mine records from a single web page by identifying similar data regions through DOM tree analysis; and Knowledge Vault [29], a method that combines different extraction approaches to feed a probabilistic knowledge base.

In 2016, Varlamov et. al. [30] argued that the degree of automation can no longer be the main classification criterion for the data extraction systems because unsupervised methods which were widely considered to be the state of the art when dealing with individual websites performed poorly or were innapropriate on cross site extraction tasks. The authors proposed a classification of methods by the extent of their application. The competing approaches were separated into two groups: methods for individual websites and methods that are applicable to whole application domains.

The first group contains most of the earlier approaches, including the supervised approaches: SRV ([14]), RAPIER ([8]), WHISK ([15]), WIEN ([2]) SoftMealy ([3]) and STALKER ([4]); and the unsupervised approaches: RoadRunner ([7]) and EXALG ([21]).

The second group is divided between domain specific methods and domain agnostic methods. Domain specific methods are designed for extracting data about a particular application domain across multiple websites. Domain specific methods integrate information about the particular application domain in the course of its development and thus are able to achieve superior performance in comparison to domain agnostic methods. One example is the method of comment extraction from blog posts described by Kao et. al. [31]. By incorporating multiple tecniques and domain specific features they are able to build a classifier that differentiates comment-blocks and non-comment blocks. This is only one of many methods tuned for various domains. Our name extractor method also belongs to this category.

Domain agnostic methods are the most general extraction methods. They can extract information from any application domain from multiple websites. They pose the hardest challenge because the tool must infer data relevance without any prior training in that particular application domain. Some examples are: ODE ([32]), ObjectRunner ([33]), and AMBER ([34]). These approaches are broader but yield worse results than domain specific methods.

3 NAME EXTRACTION

The name extraction problem is no different than a Named Entity Recognition problem, however approaches that tipically achieve high accuracy on NER tasks like Conditional Random Fields, Hidden Markov Models or Maximum Entropy Models do not necessarily perform so well when handling tabular data, as it is most common on web data extraction tasks. For example, in narrative text we could learn patterns such as "X spoke to Y" and then discover from the sentence "Alice spoke to Bob" that Alice and Bob are names.

However in the name extraction task, the lack of context words prevents these methods from detecting useful patterns over sequences of tokens. Instead we rely on a priori probabilities over

In the name extraction problem, structural HTML features and proper noun conditional probabilities play a key role.

Let $t = (t_1, t_2, ..., t_n)$ be a sequence of tokens, and $y = (y_1, y_2, ..., y_n)$ be a sequence of labels where $y_i \in \{N, W\} \forall i \in \mathbb{N}$, such that $y_i = N$ means token i is a name and $y_i = W$ means token i is a word. The problem of extracting names from a sequence of tokens is equivalent to the problem of finding an optimal sequence of labels y^* for a sequence of tokens t:

$$y^* = \underset{y}{\arg\max} \, P(t_i = y_i, t_{i+1} = y_{i+1}, ..., t_n = y_n) \qquad (1)$$

where $P(t_i = y_i)$ is the probability that token t_i has label y_i. We may employ the chain rule to explore the relationship between the joint and conditional probabilities. Consider that $P(Y_i) = P(t_i = y_i)$

$$P(Y_1, Y_2, ..., Y_n) = P(Y_n)P(Y_2|Y_1)...P(Y_n|Y_1, Y_2, ..., Y_{n-1}) \quad (2)$$

we could approximate these conditional probabilities the same way we would do for a n-gram language model. However, the conditional probabilities $P(t_i = y_i | t_{i-1} y_{i-1}, ...)$ are hard to estimate because the joint distribution $P(t_i, y_i)$ depends both on the previous label and the previous token. If we express them in terms of joint probabilities the problem becomes clearer:

$$P(t_i, y_i | t_{i-1}, y_{i-1}, t_{i-2}, y_{i-2}, ...) \qquad (3)$$

so we make the assumption that the probability that token t_i has label y_i depends on the values of previous labels but is independent of the previous tokens. For example, given a sequence of tokens $\{John, Hall\}$, the conditional probability $P(Hall|John = name)$ is equivalent to $P(Hall|any \, name)$. In other words, this means that the probability of Hall being a last name is the same regardless of a person's first name, as long as we can make sure that the previous token is a name. The conditional probabilities then become:

$$P(t_i, y_i | t_{i-1}, y_{i-1}, t_{i-2}, y_{i-2}, ...) = P(t_i, y_i | y_{i-1}, y_{i-2}, ...) \quad (4)$$

Replacing equation 5 in equation 2, yields:

$$P(Y_1, Y_2, \ldots, Y_n) = P(Y_1)P(Y_2|y_1) \ldots P(Y_n|y_1, y_2, ..., y_{n-1}) \quad (5)$$

We can once again employ the chain rule of probability to write:

$$P(x_i, y_i | y_1, y_2, \ldots, y_{i-1}) = P(y_i | y_1, y_2, \ldots, y_{i-1})P(x_i | y_1, y_2, \ldots, y_i) \quad (6)$$

Approximating $P(x_i | y_1, y_2, ..., y_i)$ by $P(x_i | y_i)$ we obtain:

$$P(x_i, y_i | y_1, y_2, \ldots, y_{i-1}) = P(y_i | y_1, y_2, \ldots, y_{i-1})P(x_i | y_i) \quad (7)$$

Take notice that $P(Y_i)$ is a different notation for $P(x_i = y_i)$ or $P(x_i, y_i)$. So, by replacing equation 7 in equation 5 we finally obtain:

$$P(Y_1, Y_2, \ldots, Y_n) = P(y_1, y_2, \ldots, y_n)P(x_1|y_1)P(x_2|y_2) \ldots P(x_n|y_n) \quad (8)$$

Equation 8 can be split into two parts: the prior, given by the first part of the equation on the right side and the conditional token probabilities, given by the rest of the equation on the right side.

3.0.1 Prior probabilities. In order to obtain the prior probability, we need to acquire estimates for all possible sequences of labels. Considering that label y_i must be either a name or a word, then there are 2^k different combinations for a window of size k.

Let N be a name label and W be a word label, then for a window of size k we would need to estimate prior probabilities for all 2^k possible sequences of labels. In practice, a window of size 4 seems to be accurate enough. In this case, the sequences would be: $\{W, W, W, W\}, \{W, W, W, N\}, \{W, W, N, W\}, \ldots$.

When names occur next to each other we have no way to tell where the first name ends and the second one starts. In order to delimit name boundaries we need to estimate priors for different sequences of name labels in addition to our previous priors. Let the first and second name labels be N_1 and N_2, respectively. Then, we need to estimate priors for the sequences: $\{W, N_1, N_1, N_2\}$, $\{W, N_1, N_1, N_2\}, \{W, N_1, N_2, N_2\}, \ldots$. In practice, we are never interested in isolated occurrences of name labels so we can exclude combinations such as $\{W_1, N_1, N_2, N_2\}$.

Most of the times when names happen inside a list they tend to be contained inside a single HTML element. Eventhough this is not always the case, this knowledge can be incorporated as an additional piece of evidence in our model. This evidence becomes specially useful when we are trying to delimit name boundaries. Let $*$ indicate a breaking point in a sequence of labels, so $W, W*, N, N$ means that the tokens taking labels WW are contained inside a single HTML element, while the remaining tokens are inside different HTML elements. We could estimate sequences with multiple breaking points, however a single breaking point has shown good results. For our window of size 4, we need to estimate all prior probabilities with the 4 possible breaking points: $\{y_1, y_2, y_3, y_4\}, \{y_1*, y_2, y_3, y_4\}$, $\{y_1, y_2*, y_3, y_4\}, \{y_1, y_2, y_3*, y_4\}$.

3.0.2 Conditional token probabilities. We need to estimate conditional token probabilities for both labels: names and words. So we need to know $P(t_i|N)$, the probability that a name is t_i and $P(t_i|W)$, the probability that a word is t_i.

For our experiments, the conditional token probabilities were obtained by maximum likelihood estimation with Laplace smoothing to account for tokens that didn't occur in the corpus. The $P(t_i|N)$ probabilities were estimated over a collection of approximately 1.5 million names from the DBLP database. The $P(t_i|W)$ probabilities were estimated over a corpus of 100 thousand documents obtained during the crawling stage. In the latter case all capitalized words were ignored when estimating probabilities in order to remove most names from the corpus.

Token probabilities can be made more precise by incorporating features in equation 8. We do that by changing the token conditional probabilities to:

$$P(t_i, f_1, f_2, \ldots, f_n|y_i) = P(t_i|y_i)P(f_1|y_i) \ldots P(f_n|y_i) \qquad (9)$$

where f_i are features, which are assumed to be independent between themselves. The features can be textual or structural. Textual features take textual clues like previous words and token length to predict if a given token is a name. Structural features infer token probabilities based on the HTML structure like tag names and nesting depth. In practice previous and next words weren't particularly effective in empirical tests, possibly due to the fact that most names occured in lists in the test collection. So they were removed from the features list.

Structural features estimated over the entire corpus end up being too general so they help little in increasing token probability estimates. HTML structure varies radically between different documents such that the only stable characteristic is that names tend to have similar structural contexts in the same faculty directory. For example, if all names appear inside a ¡tr¿ tag in a given document it does not mean that names tend to appear inside ¡tr¿ rather than any other tag in other documents. However for that particular document we may be able to identify other names and exclude words by knowing that tokens occuring inside ¡tr¿ tags have a higher probability of being names.

If our basic algorithm (without structural features) was able to extract a good number of names from a page on the first passing, we may use their structural context to estimate probabilities for our structural features. Of course those tokens that were tagged as words can be used to estimate the word conditional probabilities. On a second passing we can incorporate these improved estimates to boost the model's precision considerably. This process can be repeated multiple times to further increase peformance.

Id	Feature	Description
1	Token incidence	How often a token happens in a document
2	Token length	The token's character length
3	First+Second parents	First and second HTML parents combined
4	Third parent	Third HTML parent
5	CSS class name	Innermost CSS class valid for the token
6	Child number	The child number relative to the HTML parent
7	Nesting depth	The number of HTML parents up to root

Table 1. Features

Table 1 describes the features tested in our experiments. A couple of other features were tested, but they weren't included in the analysis due to the lack of significant impact on the results. Some of them were: next word indicates some location (street, avenue, etc.), previous word is an honorific, token is capitalized, token is the first child in an HTML element, etc.

3.1 Experiments

Features	Precision	Recall	F-measure
None	91.71% -	91.39% -	0.9155
1	91.28% (0.3880)	93.79% (0.1836)	0.9252 (0.3026)
2	92.00% (0.4189)	92.03% (0.4086)	0.9201 (0.4042)
1+2	91.83% (0.4660)	93.90% (0.1586)	0.9285 (0.2336)

Table 2. Textual features experiment

The test collection was a set of 310 manually labeled faculty directory pages. For each model the precision, recall and f-measures were calculated. All measures were tested for statistical significance with a one tailed paired T-test, the t-values are presented inside parenthesis next to each measure. Results were considered to be significant when the t-value was smaller than 0.05.

Extracted names were only considered to be correct when they resulted on an exact match with the test data. All measures were obtained by the averaged results of a 5 fold cross validation run. We first compared the textual features on our base model (NE), which passes only one time over the token list. The results are presented in table 2.

Features	Precision	Recall	F-measure
3	93.43% (0.1280)	**96.14% (0.361)**	0.9477 (0.475)
4	93.04% (0.1965)	95.67% (0.508)	0.9433 (0.740)
5	93.04% (0.1781)	95.50% (0.653)	0.9425 (0.819)
6	92.70% (0.2307)	**95.60% (0.451)**	0.9413 (0.725)
7	92.83% (0.2374)	**95.79% (0.467)**	0.9429 (0.789)
3 + 5 + 7	94.23% (0.0664)	**97.14% (0.163)**	**0.9566 (0.212)**

Table 3. Structural features experiment

Next, we tested the structural features in a model that passes two times over the token list, estimating the structural feature probabilities on the second passing. We used the best set of textual features found in the previous experiment. The results are presented in table 3.

# passings	Precision	Recall	F-measure
2	94.40% (0.525)	**97.48% (0.117)**	**0.9592 (0.151)**
3	94.37% (0.544)	**97.58% (0.104)**	**0.9595 (0.141)**
4	94.40% (0.523)	**97.61% (0.100)**	**0.9597 (0.135)**
5	94.37% (0.544)	**97.61% (0.100)**	**0.9596 (0.138)**

Table 4. Multiple passings experiment

Finally we compared models with 1, 2, 3, 4 and 5 passings over the token list using the best set of features from the previous experiments. The results are presented on table 4.

By comparing the results, it becomes clear that incorporating structural features with the two passing strategy improved all metrics significantly, even though the precision improvements came slightly above our significance threshold. By design, the features were chosen with improvements to recall rather than precision in mind. Because when a data extractor misses useful results, it becomes hard to obtain that data, however cleaning bad results is a much simpler task to be carried manually.

3.2 Further Research

Our model is finely tuned for the particular problem of researcher name extraction, but we believe this result can be generalized to solve other data extraction problems. Furthermore, the researcher affiliation data collected by the name extractor still needs to be used for evaluating the reputation metric proposed by [1] in the research group classification task.

REFERENCES

[1] Sabir Ribas, Berthier Ribeiro-Neto, Rodrygo Santos, Edmundo Souza e Silva, Alberto Ueda, and Nivio Ziviani. *Random Walks on the Reputation Graph.* ACM Press, New York, New York, USA, 2015.

[2] Nicholas Kushmerick. Wrapper induction: efficiency and expressiveness. *Artificial Intelligence*, 118(1-2):15–68, 2000.

[3] Chun Nan Hsu and Ming Tzung Dung. Generating finite-state transducers for semi-structured data extraction from the Web. *Information Systems*, 23(8):521–538, 1998.

[4] Ion Muslea, Steve Minton, and Craig Knoblock. A Hierarchical Approach to Wrapper Induction. *Proc. of the Third Annual Conf. on Autonomous Agents, ACM*, pages 190–197, 1999.

[5] Brad Adelberg. NoDoSE—a tool for semi-automatically extracting structured and semistructured data from text documents. *ACM SIGMOD Record*, 27(2):283–294, 1998.

[6] Alberto H F Laender, Berthier Ribeiro-Neto, and Altigran S da Silva. DEByE - Date extraction by example. *Data Knowl. Eng.*, 40(2):121–154, 2002.

[7] Valter Crescenzi, Giansalvatore Mecca, and Paolo Merialdo. Roadrunner: Towards automatic data extraction from large web sites. *Proceedings of the 27th International Conference on Very Large Data Bases*, pages 109–118, 2001.

[8] Mary Elaine Califf and Raymond J Mooney. Relational learning of pattern-match rules for information extraction. *Computational Linguistics*, 4:9–15, 1999.

[9] A.H.F. Laender, B. A. Ribeiro-Neto, and Juliana S.Teixeira. A brief survey of web data extraction tools. *ACM SIGMOD Record 31(2)*, pages 84–93, 2002.

[10] Joachim Hammer, Jason Mchugh, and Hector Garcia-molina. Semistructured Data : The TSIMMIS Experience. *Proceedings of the First East-European Symposium on Advances in Databases and Information Systems*, pages 1–8, 1997.

[11] Gustavo O. Arocena and Alberto O. Mendelzon. WebOQL: Restructuring documents, databases, and webs. *Theory and Practice of Object Systems*, 5(3):127–141, 1999.

[12] L. Liu, C. Pu, and W. Han. XWRAP: an XML-enabled wrapper construction system for Web information sources. *Proceedings of 16th International Conference on Data Engineering*, pages 611–621, 2000.

[13] Arnaud Sahuguet and Fabien Azavant. Building light-weight wrappers for legacy Web data-sources using W4F. *Proceedings of the 25th VLDB Conference*, 99:738–741, 1999.

[14] Dayne Freitag. Information Extraction from HTML: Application of a General Machine Learning Approach. *Proceedings of the Fifteenth National/Tenth Conference on Artificial Intelligence/Innovative Applications of Artificial Intelligence*, pages 517–523, 1998.

[15] Stephen Soderland. Learning Information Extraction Rules for Semi-Structured and Free Text. *Machine Learning*, 34(1):233–272, 1999.

[16] Chia-Hui Chang, Mohammed Kayed, Moheb Ramzy Girgis, and Khaled F Shaalan. A Survey of Web Information Extraction Systems. *IEEE Transactions on Knowledge and Data Engineering*, 18(10):1411–1428, 2006.

[17] Andrew Hogue and David Karger. Thresher : Automating the Unwrapping of Semantic Content from the World Wide Web. *WWW '05: Proceedings of the 14th international conference on World Wide Web*, pages 86–95, 2005.

[18] C.H. Chang, C.H. Chang, S.C. Lui, and S.C. Lui. IEPAD: information extraction based on pattern discovery. *Proceedings of the 10th international conference on World Wide Web*, pages 681–688, 2001.

[19] Chia Hui Chang and Shih Chien Kuo. OLERA: Semisupervised Web-data extraction with visual support. *IEEE Intelligent Systems*, 19(6):56–64, 2004.

[20] Jiying Wang and Fred H. Lochovsky. Data extraction and label assignment for web databases. *Proceedings of the twelfth international conference on World Wide Web - WWW '03*, page 187, 2003.

[21] A. Arasu, H. Garcia-Molina, Arvind Arasu, and Hector Garcia-Molina. Extracting structured data from Web pages. *2003 ACM SIGMOD International Conference on Management of Data*, pages 337 – 348, 2003.

[22] Yanhong Zhai and Bing Liu. Web data extraction based on partial tree alignment. *Proceedings of the 14th international conference on World Wide Web - WWW '05*, page 76, 2005.

[23] Sunita Sarawagi. Information extraction. *Foundations and Trends in Databases*, 1(3):261–377, 2008.

[24] Emilio Ferrara, Pasquale De Meo, Giacomo Fiumara, and Robert Baumgartner. Web data extraction, applications and techniques: A survey. *Knowledge-Based Systems*, 70:301–323, 2014.

[25] Andreas Schulz, Jörg Lässig, and Martin Gaedke. Practical web data extraction: Are we there yet? — A short survey. *IEEE/WIC/ACM International Conference on Web Intelligence (WI), 2016*, pages 562–-567, 2016.

[26] Bernhard Krüpl, Marcus Herzog, and Wolfgang Gatterbauer. Using visual cues for extraction of tabular data from arbitrary HTML documents. *Special interest tracks and posters of the 14th international conference on World Wide Web - WWW '05*, pages 1000–-1001, 2005.

[27] Emilio Ferrara and Robert Baumgartner. Automatic wrapper adaptation by tree edit distance matching. *Smart Innovation, Systems and Technologies*, 8:41–54, 2011.

[28] Shengsheng Shi, Chengfei Liu, Yi Shen, Chunfeng Yuan, and Yihua Huang. AutoRM: An effective approach for automatic Web data record mining. *Knowledge-Based Systems*, 89:314–331, 2015.

[29] Xin Dong, Evgeniy Gabrilovich, Geremy Heitz, Wilko Horn, Ni Lao, Kevin Murphy, Thomas Strohmann, Shaohua Sun, and Wei Zhang. Knowledge vault: a web-scale approach to probabilistic knowledge fusion. *Proceedings of the 20th ACM SIGKDD international conference on Knowledge discovery and data mining - KDD '14*, pages 601–610, 2014.

[30] M. I. Varlamov and D. Yu. Turdakov. A survey of methods for the extraction of information from Web resources. *Programming and Computer Software*, 42(5):279–291, 2016.

[31] Huan-An Kao and Hsin-Hsi Chen. Comment Extraction from Blog Posts and Its Applications to Opinion Mining. *Proceedings of the 7th International Conference on Language Resources and Evaluation*, pages 1113–1120, 2010.

[32] Weifeng Su and Frederick H Lochvsky. ODE : Ontology-assisted Data Extraction. 1(212), 2009.

[33] Talel Abdessalem, B Cautis, and Nora Derouiche. ObjectRunner: lightweight, targeted extraction and querying of structured web data. *Proceedings of the VLDB …*, 3(2):1585–1588, 2010.

[34] Tim Furche, Georg Gottlob, Giovanni Grasso, Giorgio Orsi, Christian Schallhart, and Cheng Wang. AMBER: Automatic Supervision for Multi-Attribute Extraction. *arXiv preprint*, 1210(5984):1–22, 2012.

[35] Bing Liu and Robert Grossman. Mining data records in Web pages. *Knowledge discovery and data mining*, page 601, 2003.

[36] Eugene Garfield. Citation Indexes for Science. *Science*, 122:108–111, 1955.

[37] Nicholas Kushmerick, Daniel S Weld, and R Doorenbos. Wrapper induction for information extraction. *Intl Joint Conference on Artificial Intelligence IJCAI*, pages 729–737, 1997.

[38] Katharina Kaiser and Silvia Miksch. Information Extraction. *Technology*, (May):32, 2005.

[39] S. Flesca, G. Manco, E. Masciari, E. Rende, and A. Tagarelli. Web wrapper induction: a brief survey. *AI Communications*, 17(2):57–61, 2004.

[40] Anhai Doan, Jeffrey F Naughton, Raghu Ramakrishnan, Akanksha Baid, Xiaoyong Chai, Fei Chen, Ting Chen, Eric Chu, Pedro Derose, Byron Gao, Chaitanya Gokhale, Jiansheng Huang, Warren Shen, and Ba-quy Vuong. Information Extraction Challenges in Managing Unstructured Data. *ACM SIGMOD Record*, 37(4):14–20, 2008.

[41] Chia-hui Chang. FiVaTech2 : A Supervised Approach to Role Differentiation for Web Data Extraction From Template Pages. (March 2015), 2012.

[42] Nicholas Kushmerick and Nicholas Kushmerick. Finite-state approaches to Web information extraction. *Lecture Notes in Computer Science*, 2700:77–91, 2003.

[43] Leandro Neiva Lopes Figueiredo, Guilherme Tavares de Assis, and Anderson A. Ferreira. DERIN: A data extraction method based on rendering information and n-gram. *Information Processing & Management*, 53(5):1120–1138, 2017.

[44] Giacomo Fiumara. Automated information extraction from Web sources: A survey. *CEUR Workshop Proceedings*, 312:1–9, 2007.

[45] Salim Khalil and Mohamed Fakir. RCrawler: An R package for parallel web crawling and scraping. *SoftwareX*, 6:98–106, 2017.